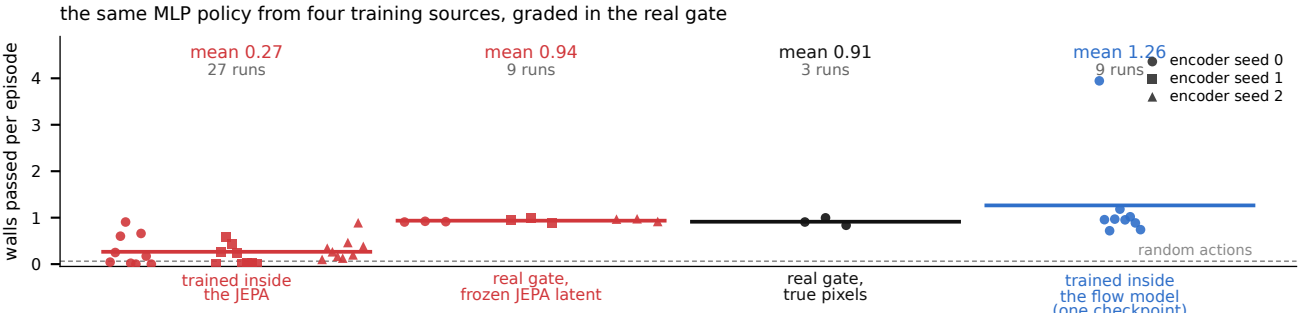
the same MLP policy from four training sources, graded in the real gate
walls passed per episode
0
1
2
3
4
mean 0.27
27 runs
mean 0.94
9 runs
mean 0.91
3 runs
mean 1.26
9 runs
encoder seed 0
encoder seed 1
encoder seed 2
random actions
trained inside
the JEPA
real gate,
frozen JEPA latent
real gate,
true pixels
trained inside
the flow model
(one checkpoint)